Udit Gupta

✉ ugupta@g.harvard.edu • ugupta.com

Education

Harvard University, Ph.D. — Cambridge, MA
Computer Science — 2016-Present
Advisors: Professor David Brooks, Professor Gu-Yeon Wei
Research Interests: Computer architecture, sustainable computing, deep learning, personalized recommendation

Harvard University, Masters of Science — Cambridge, MA
Computer Science — 2020
GPA: 3.87

Cornell University, Bachelor of Science — Ithaca, NY
Electrical & Computer Engineering, Computer Science — 2012-2016
Advisor: Professor Zhiru Zhang
GPA: 4.00, Dean's List (All semesters), *summa cum laude*

Research Experience

Harvard University — Cambridge, MA
Graduate Researcher — 2016-Present

- Detailing the environmental impact of computing across hardware life-cycles for mobile and data center scale systems.
- Accelerating DNN-based personalized recommendation with specialized schedulers, memory systems, and hardware.
- Developed benchmarks for DNN-based recommendation models based on in-depth architectural characterization.
- Designed specialized hardware to parallelize static and dynamic sparse execution in RNNs for on-chip speech recognition.
- Collaborated with graduate students and post-docs on 16nm tape-out with ARM A53 CPU and 4 coherent accelerators.

Cornell University — Ithaca, NY
Undergraduate Researcher — 2013-2016

- Developed benchmarks and optimizations for designing accelerators using high-level synthesis on FPGAs.

Industry Experience

Facebook, Inc. — Menlo Park, CA
Facebook AI Research (FAIR) Visiting Research Scientist — January 2021-Present
Facebook AI Research (FAIR) Intern — January 2020 - December 2020
AI Infrastructure Research Intern — September 2018 - January 2020

- Characterizing the architectural implications of deep learning based personalized recommendation systems.
- Designing inference schedulers and specialized hardware to optimize the performance of recommendation in datacenters under different run-time configurations such as models, server architecture, batching, and co-location.

Algo-Logic Systems — Santa Clara, CA
Hardware Design and Verification Engineering Intern — Summer 2015

- Designed and implemented OpenCL interface to software kernels with existing IP on FPGAs for financial data parsers.

Honors and Awards

IEEE MICRO Top Picks Honorable Mention — 2021
Best Paper Nominee at Parallel Architectures and Compilation Techniques (PACT) — 2019
Best Paper Nominee at Design Automation Conference (DAC) — 2018
Harvard Smith Family Fellowship — 2017
National Science Foundation GRFP Honorable Mention — 2016
Richard A. Newton Young Fellows Scholarship at *DAC 2015* — 2015

Open Source Initiatives

- DeepRecSys: A System for Optimizing End-To-End At-scale Neural Recommendation Inference (GitHub)
- Ares: A framework for quantifying the resilience of deep neural networks (GitHub)
- MLPerf: A Benchmark for Machine Learning from an Academic/Industry Cooperative (MLPerf)

Community Involvement and Professional Service

Computer Architecture Student Association (CASA)

Steering Committee Member 2020-Present

- Steering committee member for a student-run group fostering an inviting computer architecture student community.
- Organizing a summer 2021 reading group focused on advancing and promoting diversity, equity, and inclusion in computer architecture.
- Gathered qualitative and quantitative feedback on various mentoring and networking initiatives in computer architecture, published at Workshop on Computer Architecture Education (WCAE) co-located with ISCA 2021.

Workshop and Tutorial Organizing Activities

- Computing Landscapes for Environmental Accountability and Responsibility (CLEAR), Co-founder ISCA 2021
- Negative Outcomes Post-Mortems and Experiences (NOPE), Co-organizer ASPLOS 2021
- Journal of Retrospectives, Negative Results, Experiences (JOURNE), Co-founder MLSys 2021
- Personalized Recommendation Systems and Algorithms (PeRSonAl), Co-founder MLSys 2021
- Personalized Recommendation Systems and Algorithms (PeRSonAl), Co-founder ISCA 2020
- Personalized Recommendation Systems and Algorithms (PeRSonAl), Co-founder ASPLOS 2020
- Negative Outcomes Post-Mortems and Experiences (NOPE), Co-organizer ASPLOS 2019

Professional Memberships

- 3C (Cultural Competence in Computing) Fellow 2021-2022
- Computer Architecture Student Association, Steering Committee Member 2020-Present
- Harvard SEAS Graduate Council, Member 2019-Present
- Harvard SITN Lecture Series , Lecture Facilitator and Director 2018-2019
- Cornell IEEE Student Chapter, President and Corporate Director 2013-2016

Teaching Experience

Graduate Teaching Fellow

Cambridge, MA

Harvard University 2 semesters

- CS 290: PhD Grad Cohort Research Seminar Fall 2020
- CS 141: Computing Hardware Spring 2019

Research Mentor for Undergraduate Students

Cambridge, MA

Harvard University

- Advised 3 summer undergraduate students on building recommendation training zoo.
- Advised undergraduate student on *"Quantifying the Impact of Data Encoding on DNN Fault Tolerance" (Fastpath workshop)*.
- Advised undergraduate senior thesis on *"Improving Resiliency of Deep Neural Networks for Denser eNVM Storage"*.
- Mentored 2 summer undergraduate students on *"Applications of Deep Neural Networks for Ultra Low Power IoT"* (ICCD 2017) .

Education Outreach for Middle School and High School Students

Cambridge, MA

Harvard University

- Taught a 1.5 hour class on *"Sustainable computing"* to 25 high school students at *Rainstorm* in Summer 2021.
- Taught a one-hour class on *"Sustainable computing"* to 30 high school students at *MIT Splash!* in Spring 2021.
- Taught 3 hour hardware engineering course, *"Computer's Don't Byte!"*, to 25 middle- and high- school students at *Cornell Splash!* in Fall 2014 and Spring 2015.

Undergraduate Teaching Assistant — Ithaca, NY
Cornell University — 4 semesters
- CS 3420/ECE 3140: Embedding Systems — Spring 2016
- EdX MOOC: The Computing Inside Your Smartphone — Summer 2014
- ECE 2300: Introduction to Digital Logic and Computer Organization — Spring 2014, Fall 2015, Spring 2015

Publications

RecPipe: Co-designing Models and Hardware to Jointly Optimize Recommendation Quality and Performance
Udit Gupta, Samuel Hsia, Jeff Zhang, Mark Wilkening, Javin Pombra, Hsien-Hsin S. Lee, Gu-Yeon Wei, Carole-Jean Wu, David Brooks
IEEE/ACM International Symposium on Microarchitecture (MICRO 2022).

RecSSD: Near Data Processing forSolid State Drive Based Recommendation Inference
Mark Wilkening, **Udit Gupta**, Samuel Hsia, Caroline Trippel, Carole-Jean Wu, David Brooks, Gu-Yeon Wei
International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS 2021)

Chasing Carbon: The Elusive Environmental Footprint of Computing
Udit Gupta, Young Geun Kim, Sylvia Lee, Jordan Tse, Hsien-Hsin S. Lee, Gu-Yeon Wei, David Brooks, Carole-Jean Wu
IEEE International Symposium on High-Performance Computer Architecture (HPCA 2021)
Featured by: Harvard Gazette, Tech at Facebook, Bloomberg Green

Cross-Stack Workload Characterization of Deep Recommendation Systems
Samuel Hsia, **Udit Gupta**, Mark Wilkening, Carole-Jean Wu, Gu-Yeon Wei, David Brooks
IEEE International Symposium on Workload Characterization (IISWC 2020)

DeepRecSys: A System for Optimizing End-To-End At-scale Neural Recommendation Inference
Udit Gupta, Samuel Hsia, Vikram Saraph, Xiaodong Wang, Brandon Reagen, Gu-Yeon Wei, Hsien-Hsin S. Lee, David Brooks, Carole-Jean Wu
The 47th IEEE/ACM International Symposium on Computer Architecture (ISCA 2020).

RecNMP: Accelerating Personalized Recommendation with Near-Memory Processing
Liu Ke, **Udit Gupta**, Carole-Jean Wu, Benjamin Youngjae Cho, Mark Hempstead, Brandon Reagen, Xuan Zhang, David Brooks, Vikas Chandra, Utku Diril, Amin Firoozshahian, Kim Hazelwood, Bill Jia, Hsien-Hsin S. Lee, Meng Li, Bert Maher, Dheevatsa Mudigere, Maxim Naumov, Martin Schatz, Mikhail Smelyanskiy, Xiaodong Wang
The 47th IEEE/ACM International Symposium on Computer Architecture (ISCA 2020).

Architectural Implications of Facebook's DNN-based Personalized Recommendation
Udit Gupta, Xiaodong Wang, Maxim Naumov, Carole-Jean Wu, Brandon Reagen, David Brooks, Bradford Cottel, Kim Hazelwood, Bill Jia, Hsien-Hsin S. Lee, Andrey Malevich, Dheevatsa Mudigere, Mikhail Smelyanskiy, Liang Xiong, Xuan Zhang
IEEE International Symposium on High-Performance Computer Architecture (HPCA 2020)
★ ***IEEE MICRO Top Picks 2020 - Honorable Mention*** ★
Featured by: Facebook Research

MASR: A Modular Accelerator for Sparse RNNs
Udit Gupta, Brandon Reagen, Lillian Pentecost, Marco Donato, Thierry Tambe, Alexander Rush, Gu-Yeon Wei, David Brooks
Parallel Architectures and Compilation Techniques (PACT 2019).
★ ***Best Paper Nominee*** ★

MaxNVM: Maximizing DNN Storage Density and Inference Efficiency with Sparse Encoding and Error Mitigation
Lillian Pentecost, Marco Donato, Brandon Reagen, **Udit Gupta**, Siming Ma, Gu-Yeon Wei, David Brooks.
IEEE/ACM International Symposium on Microarchitecture (MICRO 2019).

A 16nm 25mm^2 SoC with a 54.5× Flexibility-Efficiency Range from Dual-Core Arm Cortex-A53, to eFPGA, and Cache-Coherent Accelerators
Paul Whatmough, Sae Kyu Lee, Marco Donato, Hsea-Ching Hseuh, Sam Xi, **Udit Gupta**, Lillian Pentecost, Glenn Ko, David Brooks, Gu-Yeon Wei.
Symposia on VLSI Technology and Circuits. (VLSI 2019)

SMIV: A 16nm SoC with Efficient and Flexible DNN Acceleration for Intelligent IoT Devices.
Paul Whatmough, Sae Kyu Lee, Sam Xi, **Udit Gupta**, Lillian Pentecost, Marco Donato, Hsea-Ching Hseuh, David Brooks, Gu-Yeon Wei.
Hot Chips (Hot Chips 2018).

Weightless: Lossy Weight Encoding for Deep Neural Network Compression.
Brandon Reagen, **Udit Gupta**, Robert Adolf, Michael Mitzenmacher, Alexander Rush, Gu-Yeon Wei, David Brooks.
International Conference on Machine Learning (ICML 2018).

Ares: A Framework for Quantifying the Resilience of Deep Neural Networks.
Brandon Reagen, **Udit Gupta**, Lillian Pentecost, Paul Whatmough, Sae Kyu Lee, Niamh Mulholland, Gu-Yeon Wei, David Brooks.
Design Automation Conference (DAC 2018).
★ ***Best Paper Nominee*** ★

On-chip Deep Neural Network Storage with Multi-level eNVM.
Marco Donato, Brandon Reagen, Lillian Pentecost, **Udit Gupta**, David Brooks, Gu-Yeon Wei.
Design Automation Conference (DAC 2018).

Rosetta: A Realistic Benchmark Suite for Software Programmable FPGAs.
Yuan Zhou, **Udit Gupta**, Steve Dai, Ritchie Zhao, Nitish Srivastava, Hanchen Jin, Joseph Featherston, Yi-Hsiang Lai, Gai Liu, Gustavo Velasquez, Wenping Wang, Zhiru Zhang.
ACM/SIGDA International Symposium on Field-Programmable Gate Arrays (FPGA 2018)

Dynamic Hazard Resolution for Pipelining Irregular Loops in High-Level Synthesis.
Steve Dai, Ritchie Zhao, Gai Liu, Shreesha Srinath, **Udit Gupta**, Christopher Batten, Zhiru Zhang.
ACM/SIGDA International Symposium on Field-Programmable Gate Arrays (FPGA 2017)

Mapping-Aware Constrained Scheduling for LUT-Based FPGAs.
Mingxing Tan, Steve Dai, **Udit Gupta**, Zhiru Zhang.
ACM/SIGDA International Symposium on Field-Programmable Gate Arrays (FPGA 2015)

Technical Articles

Deep Learning: It's Not All About Recognizing Cats and Dogs
Carole-Jean Wu, David Brooks, **Udit Gupta** , Hsien-Hsin Lee, and Kim Hazelwood
ACM SIGARCH, Computer Architecture Today

Designing AI-Enabled Technology for Society
Udit Gupta, Lillian Pentecost
Harvard SITN, October 2018